

Lecture on

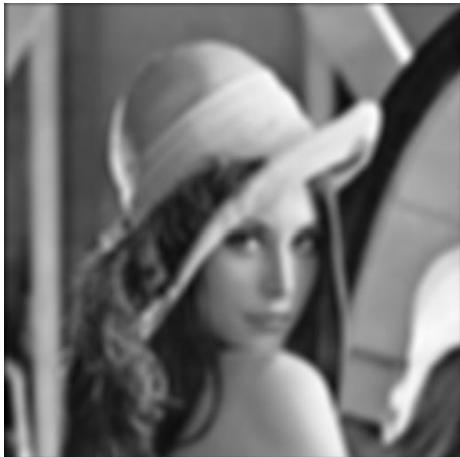
Image Filtering

Credit:
Davide Scaramuzza
University of Zurich <http://rpg.ifi.uzh.ch>

Image filtering

- The word *filter* comes from frequency-domain processing, where “filtering” refers to the process of accepting or rejecting certain frequency components
- We distinguish between low-pass and high-pass filtering
 - A **low-pass filter** smooths an image (retains low-frequency components)
 - A **high-pass filter** retains the contours (also called edges) of an image (high frequency)

Low-pass filtered image



High-pass filtered image



Low-pass filtering

Low-pass filtering

Motivation: noise reduction

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian distribution



Original



Salt and pepper noise

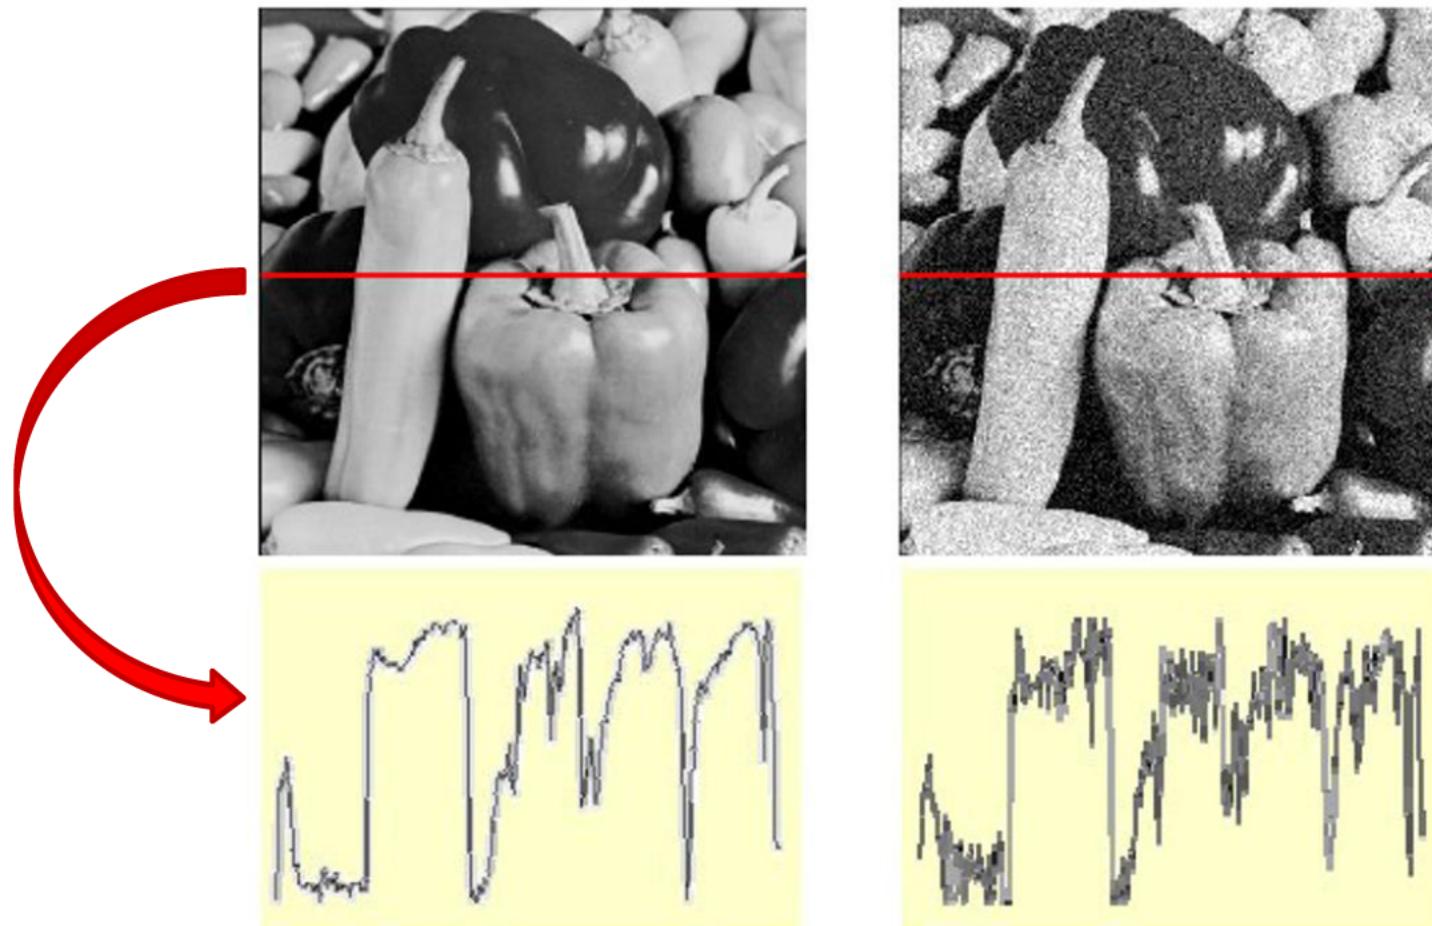


Impulse noise



Gaussian noise

Gaussian noise



$$I(x, y) = \hat{I}(x, y) + \eta(x, y)$$

$\eta(x, y) \sim N(\mu, \sigma)$

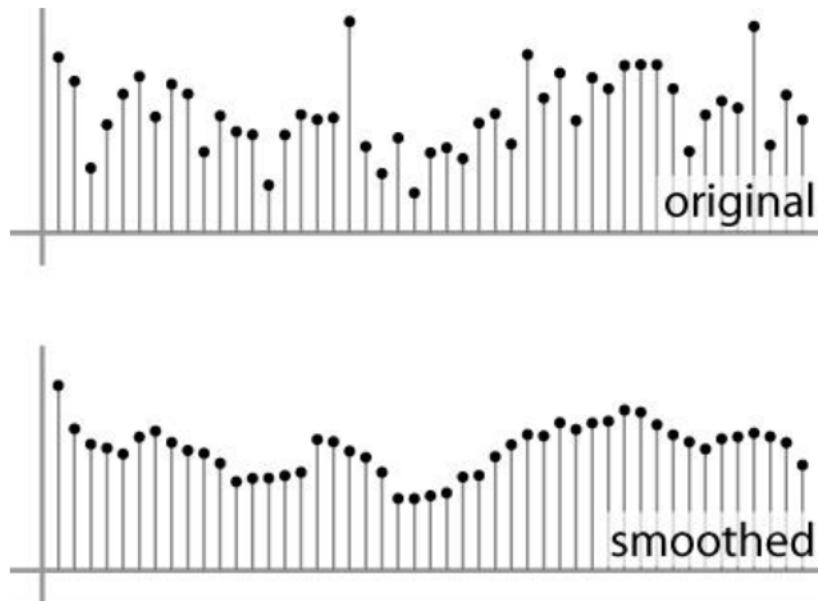
How could we reduce the noise to try to recover the “ideal image”?

Moving average

- Replaces each pixel with an average of all the values in its neighborhood
- Assumptions:
 - Expect pixels to be like their neighbors
 - Expect noise process to be independent from pixel to pixel

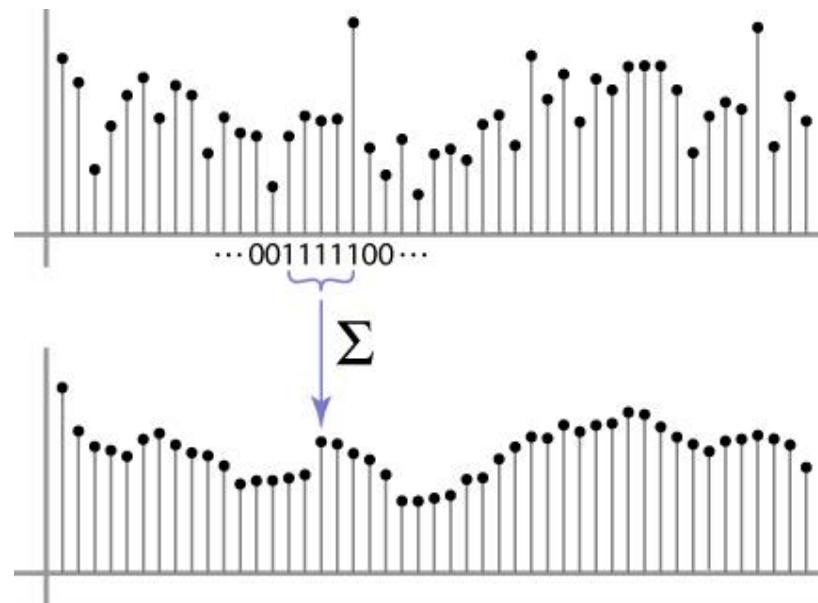
Moving average

- Replaces each pixel with an average of all the values in its neighborhood
- Moving average in 1D:



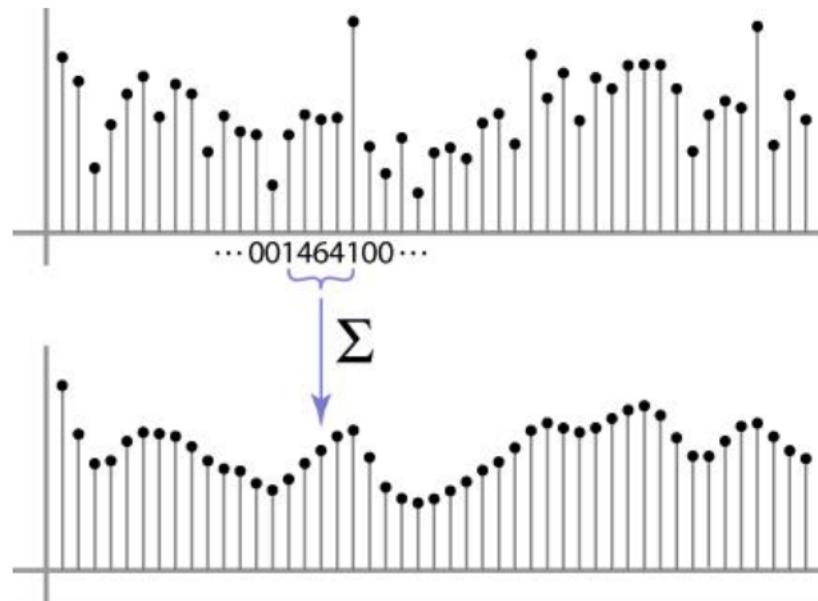
Weighted Moving Average

- Can add weights to our moving average
- *Weights* [1, 1, 1, 1, 1] / 5



Weighted Moving Average

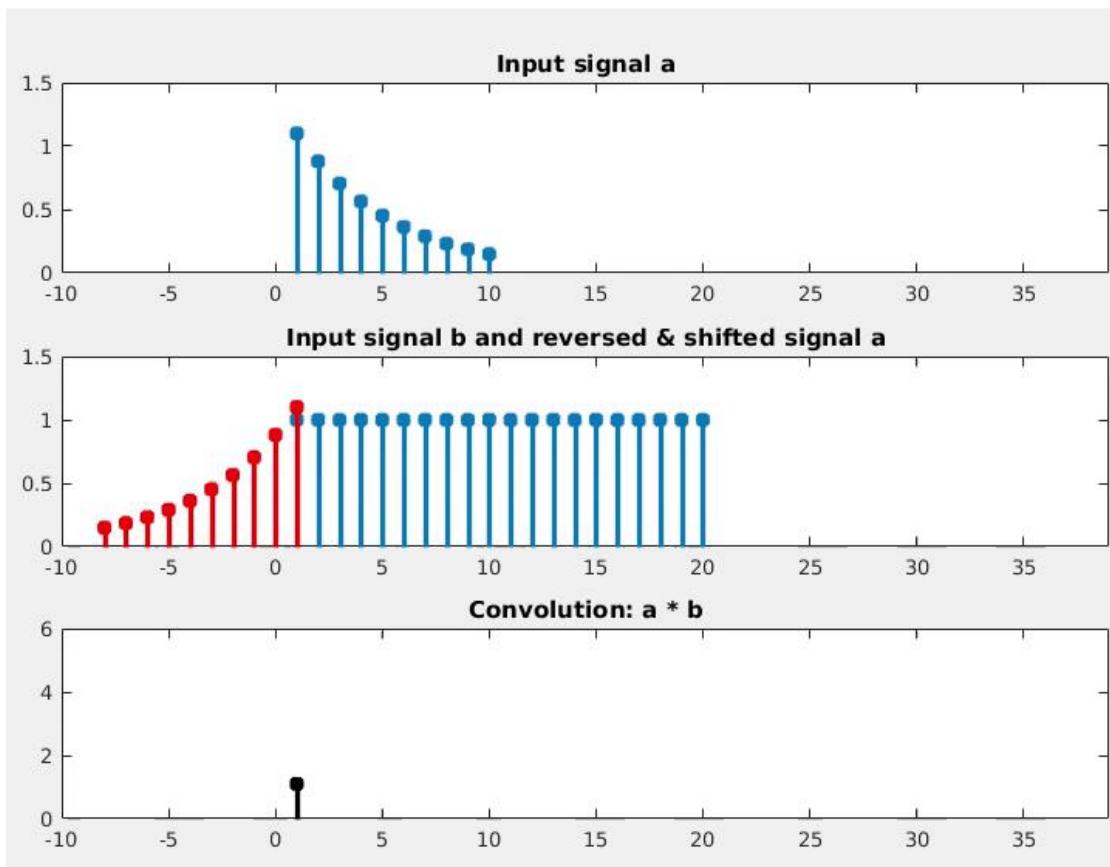
- Non-uniform weights $[1, 4, 6, 4, 1] / 16$



This operation is called *convolution*

Example of convolution of two sequences (or “signals”)

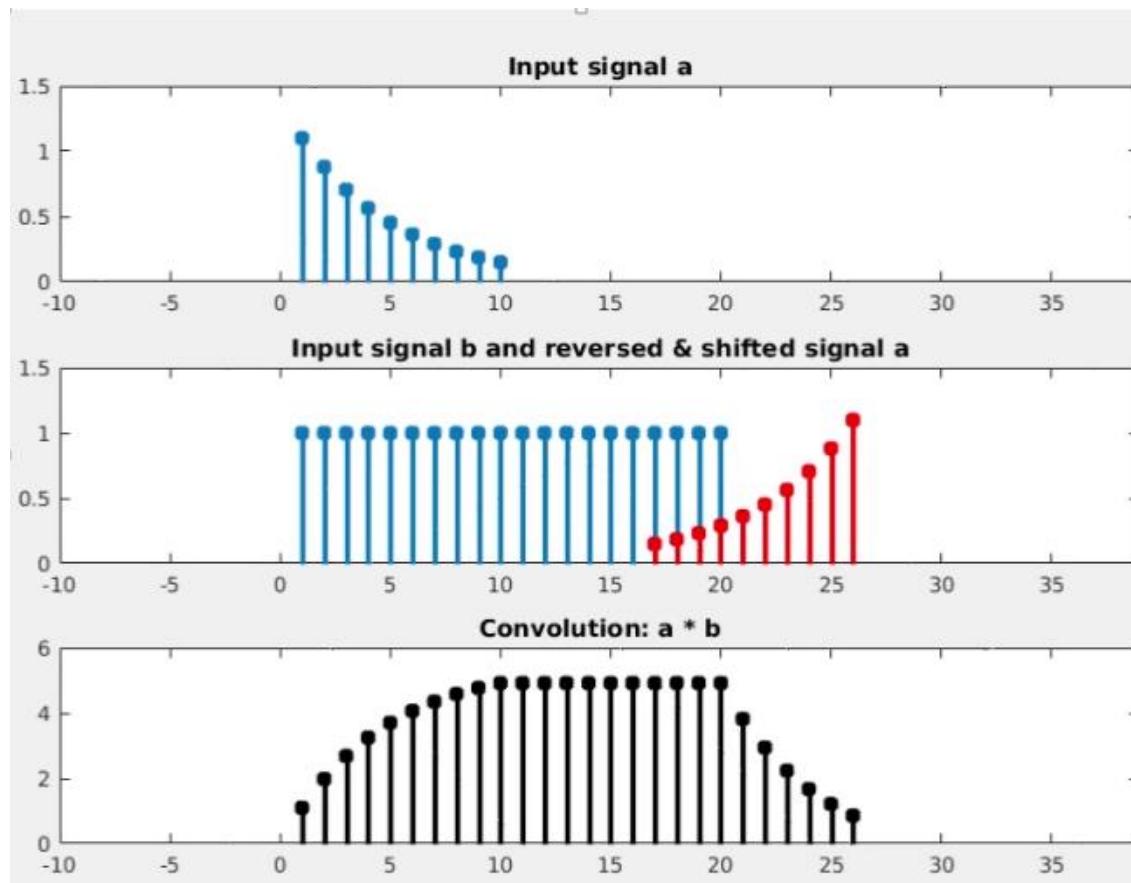
- One of the sequences is flipped (right to left) before sliding over the other
- Notation: $a * b$
- Nice properties: linearity, associativity, commutativity, etc.



This operation is called *convolution*

Example of convolution of two sequences (or “signals”)

- One of the sequences is flipped (right to left) before sliding over the other
- Notation: $a * b$
- Nice properties: linearity, associativity, commutativity, etc.

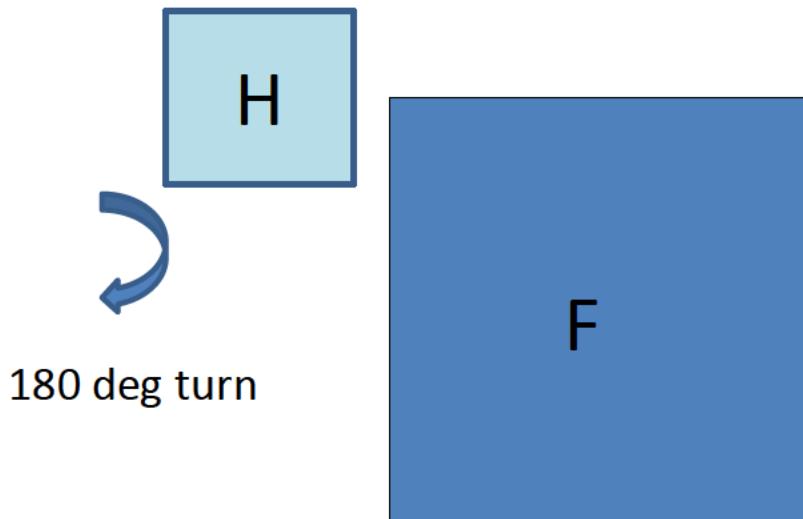


2D Filtering

- Convolution:
 - Flip the filter in both dimensions (bottom to top, right to left) (=180 deg turn)
 - Then slide the filter over the image and compute sum of products

$$G[x, y] = \sum_{u=-k}^k \sum_{v=-k}^k F[x - u, y - v] H[u, v]$$

$$G = F * H$$



Filtering an image: replace each pixel
with a linear combination of its neighbors.

The **filter H** is also called “**kernel**” or “**mask**”.

Review: Convolution vs. Cross-correlation

Convolution $G = F * H$

$$G[x, y] = \sum_{u=-k}^k \sum_{v=-k}^k F[x - u, y - v]H[u, v]$$

Properties: linearity, associativity, commutativity

Cross-correlation $G = F \otimes H$

$$G[x, y] = \sum_{u=-k}^k \sum_{v=-k}^k F[x + u, y + v]H[u, v]$$

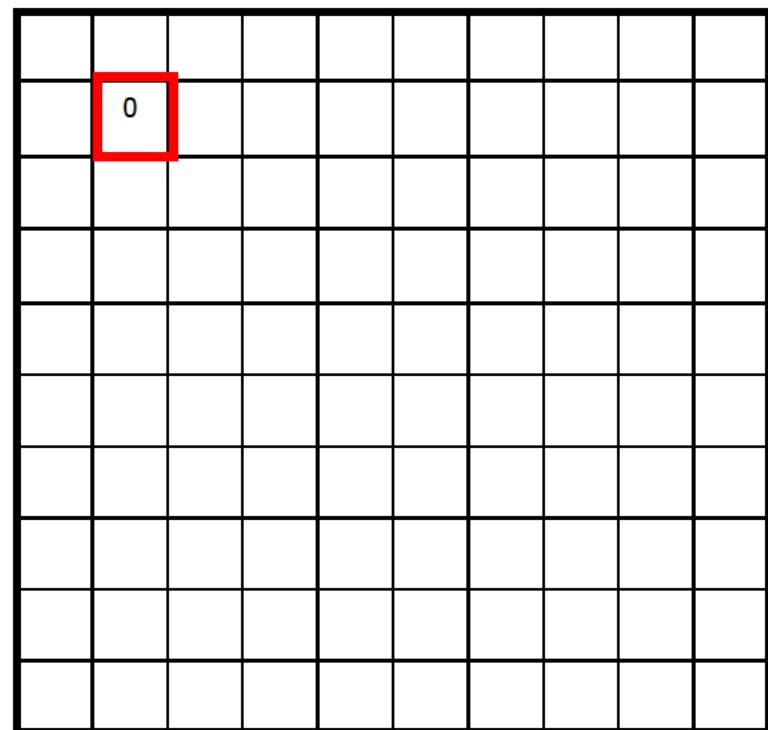
Properties: linearity, but not associativity and commutativity

For a Gaussian or box filter, will the output of convolution and correlation be different?

Example: Moving Average in 2D

Input image	$F[x, y]$
“box filter”	$F[x, y]$

Filtered image
 $G[x, y]$



Example: Moving Average in 2D

Input image

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

Filtered image

$$G[x, y]$$

0	10									

Example: Moving Average in 2D

Input image

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	90	90	0	90	90	90	0
0	0	0	90	90	90	90	90	90	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filtered image

$$G[x, y]$$

			0	10	20				

Example: Moving Average in 2D

Input image

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filtered image

$$G[x, y]$$

			0	10	20	30			

Example: Moving Average in 2D

Input image

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filtered image

$$G[x, y]$$

Example: Moving Average in 2D

Input image

$$F[x, y]$$

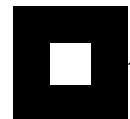
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Filtered image

$$G[x, y]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

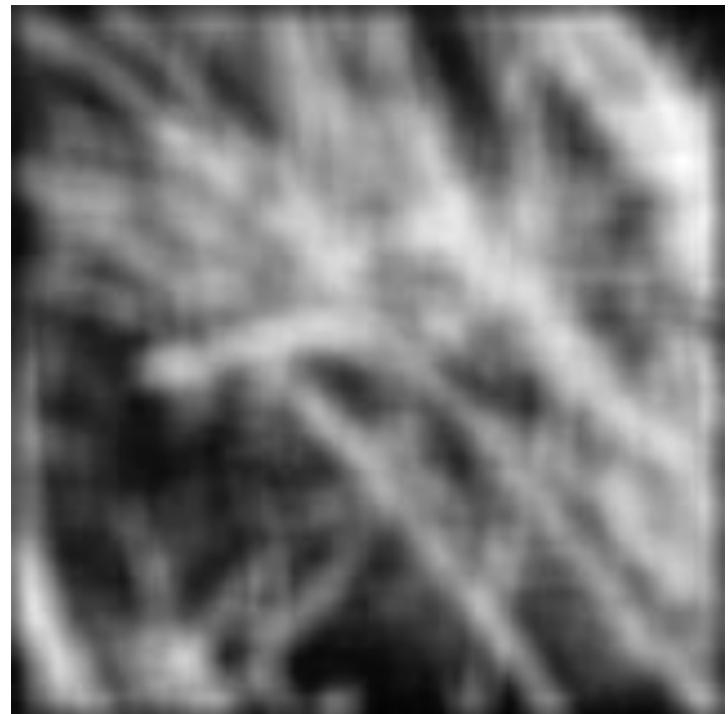
Example: Moving Average in 2D



Box filter:
white = high value, black = zero value



original



filtered

Gaussian filter

- What if we want the closest pixels to have higher influence on the output?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$F[x, y]$$

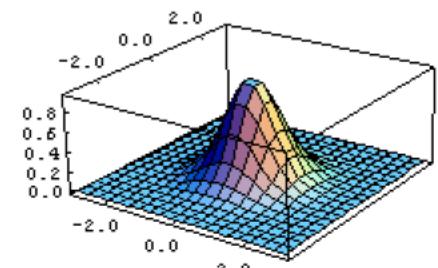
$$\frac{1}{16}$$

1	2	1
2	4	2
1	2	1

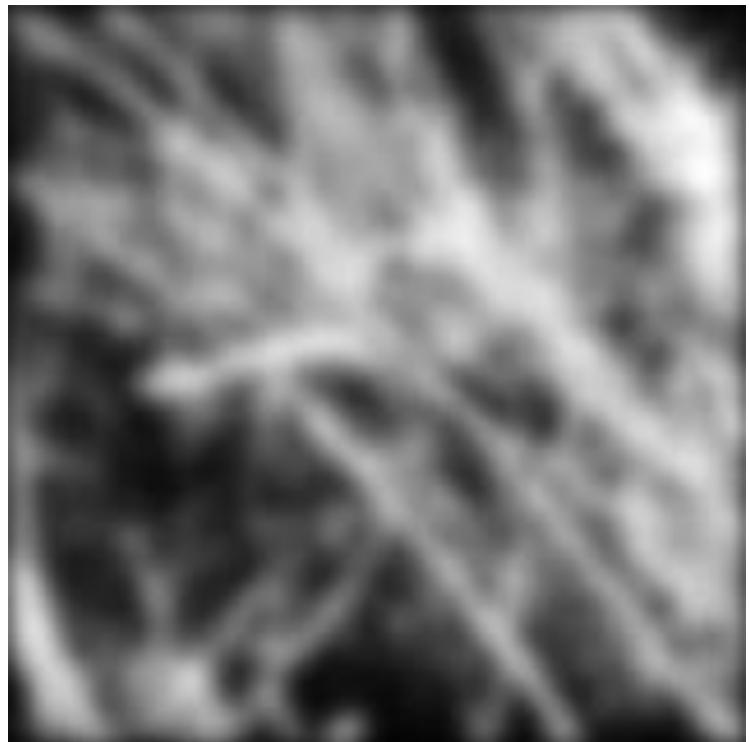
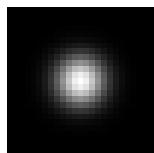
$$H[u, v]$$

This kernel is the approximation of a Gaussian function:

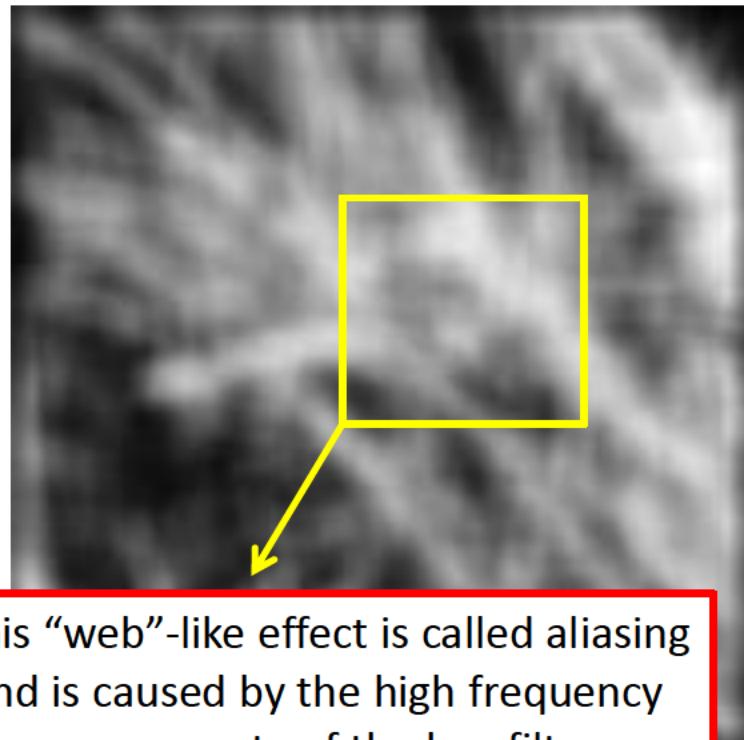
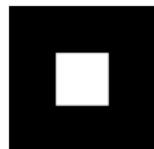
$$H[u, v] \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



Smoothing with a Gaussian



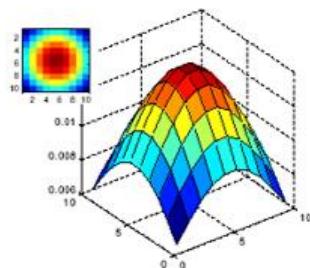
Compare the result with a box filter



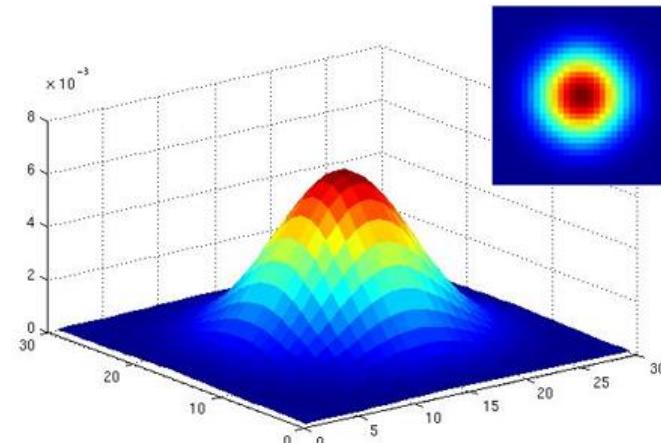
This “web”-like effect is called aliasing and is caused by the high frequency components of the box filter

Gaussian filters

- What parameters matter?
- **Size** of the kernel
 - NB: a Gaussian function has **infinite support**, but discrete filters use finite kernels



$\sigma = 5$ pixels
with 10×10 pixel kernel

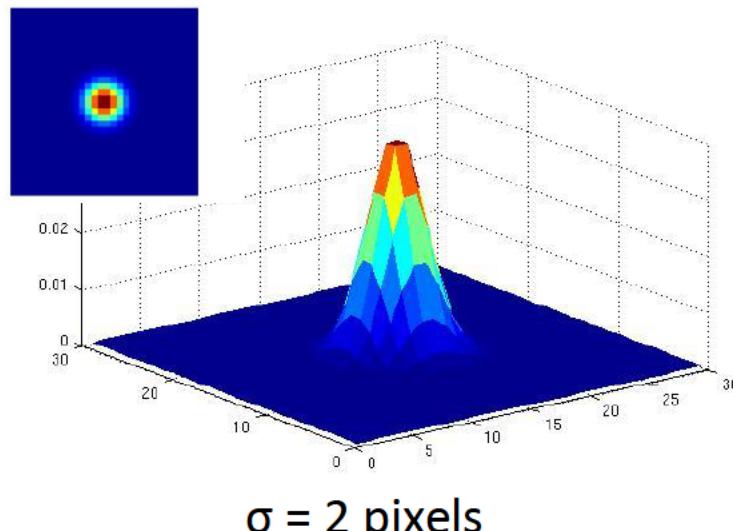


$\sigma = 5$ pixels
with 30×30 pixel kernel

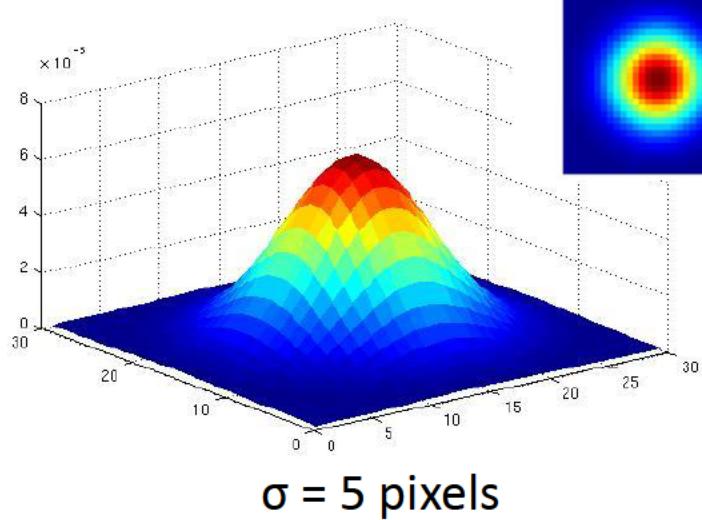
Which one approximates better the ideal Gaussian filter, the left or the right one?

Gaussian filters

- What parameters matter?
- **Variance** of Gaussian: control the amount of smoothing



$\sigma = 2$ pixels
with 30×30 pixel kernel

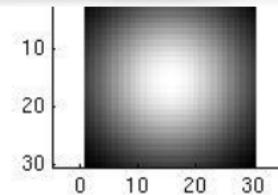
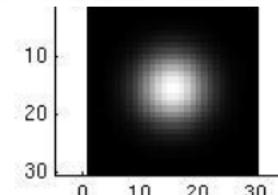
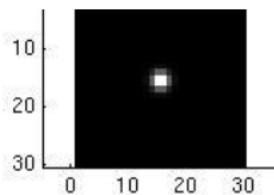


$\sigma = 5$ pixels
with 30×30 pixel kernel

Recall: standard deviation = σ [pixels], variance = σ^2 [pixels 2]

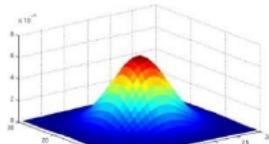
Smoothing with a Gaussian

σ is called “scale” of the Gaussian kernel, and controls the amount of smoothing.

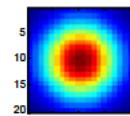


Sample Matlab code

```
>> hsize = 20;  
>> sigma = 5;  
>> h = fspecial('gaussian', hsize, sigma);
```



```
>> mesh(h);
```



```
>> im = imread('panda.jpg');  
>> outim = imfilter(im, h);  
>> imshow(outim);
```



im



outim 28

Boundary issues

- What about near the image edges?
 - the filter window falls off the edges of the image
 - need to pad the image borders
 - methods:
 - zero padding (black)
 - wrap around
 - copy edge
 - reflect across edge



Summary on (linear) smoothing filters

- Smoothing filter
 - has positive values (also called coefficients)
 - **sums to 1** → preserve brightness of constant regions
 - removes “high-frequency” components; “low-pass” filter

Non-linear filtering

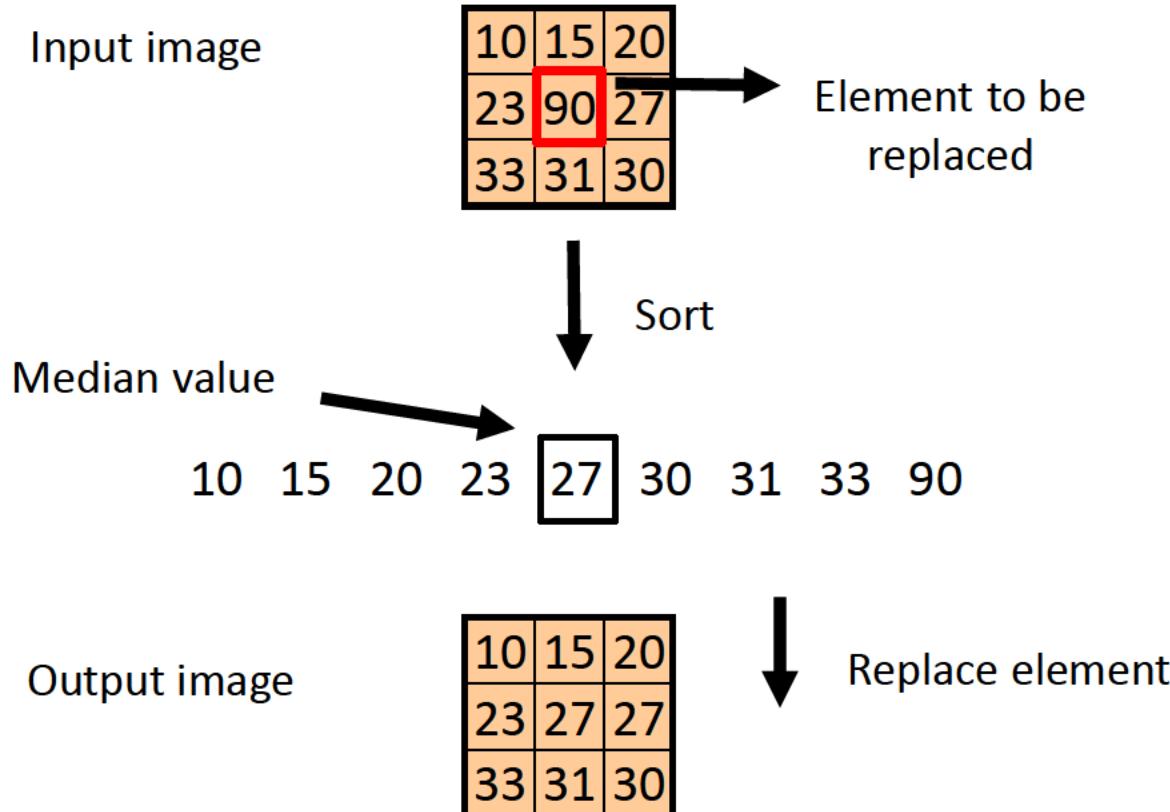
Effect of smoothing filters



Linear smoothing filters do not alleviate salt and pepper noise!

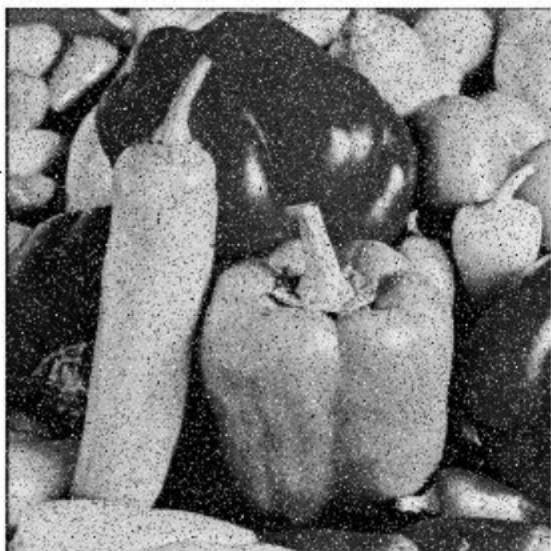
Median filter

- It is a non-linear filter
- Removes spikes: good for “*impulse noise*” and “*salt & pepper noise*”

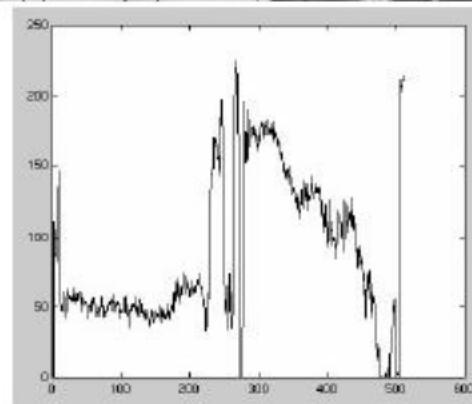
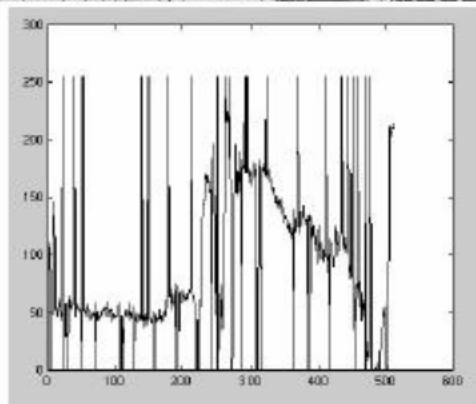


Median filter

Salt and pepper noise →



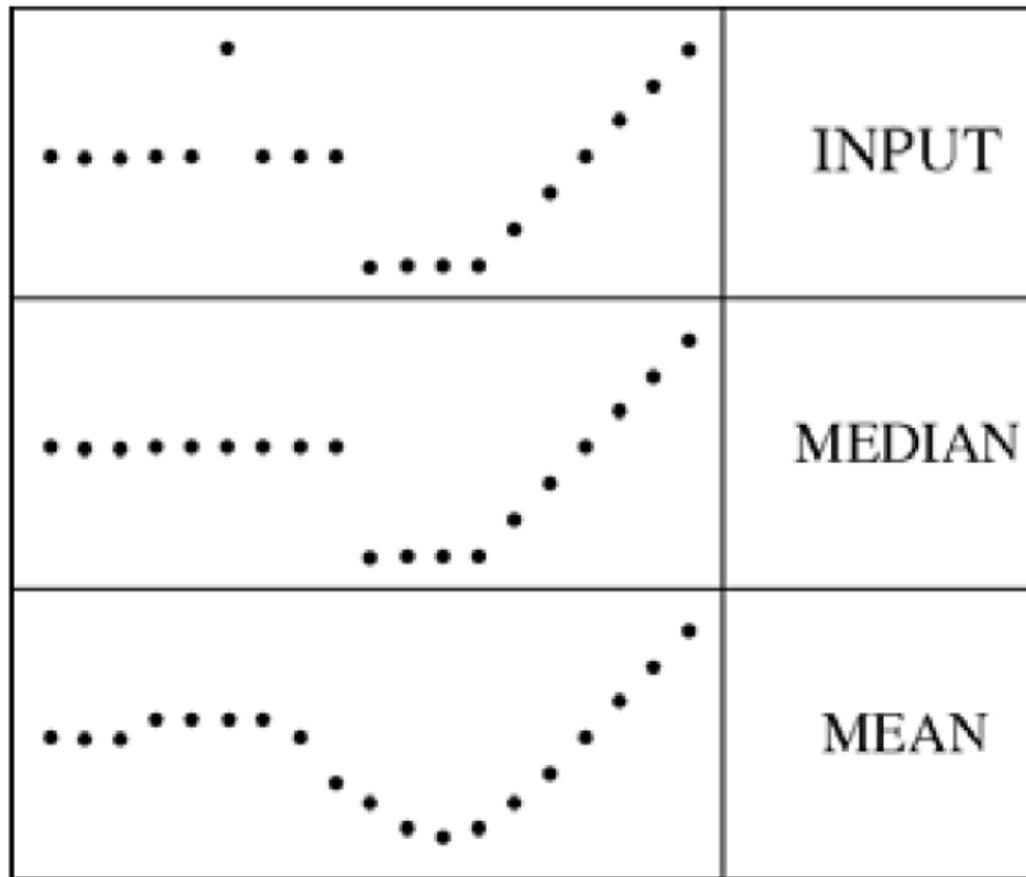
← Median filtered



Plots of a row of the image

Median filter

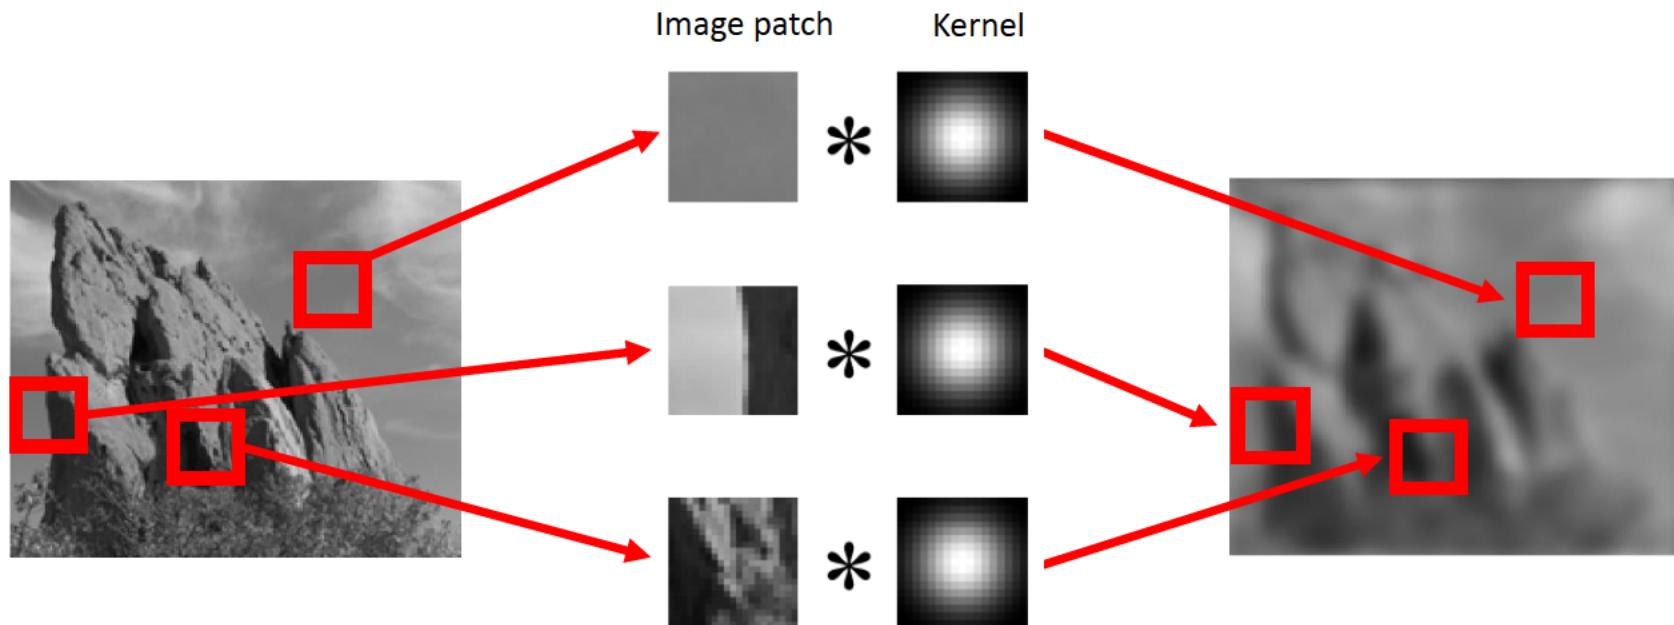
- Median filter preserves strong edges,



... but it removes small edges.

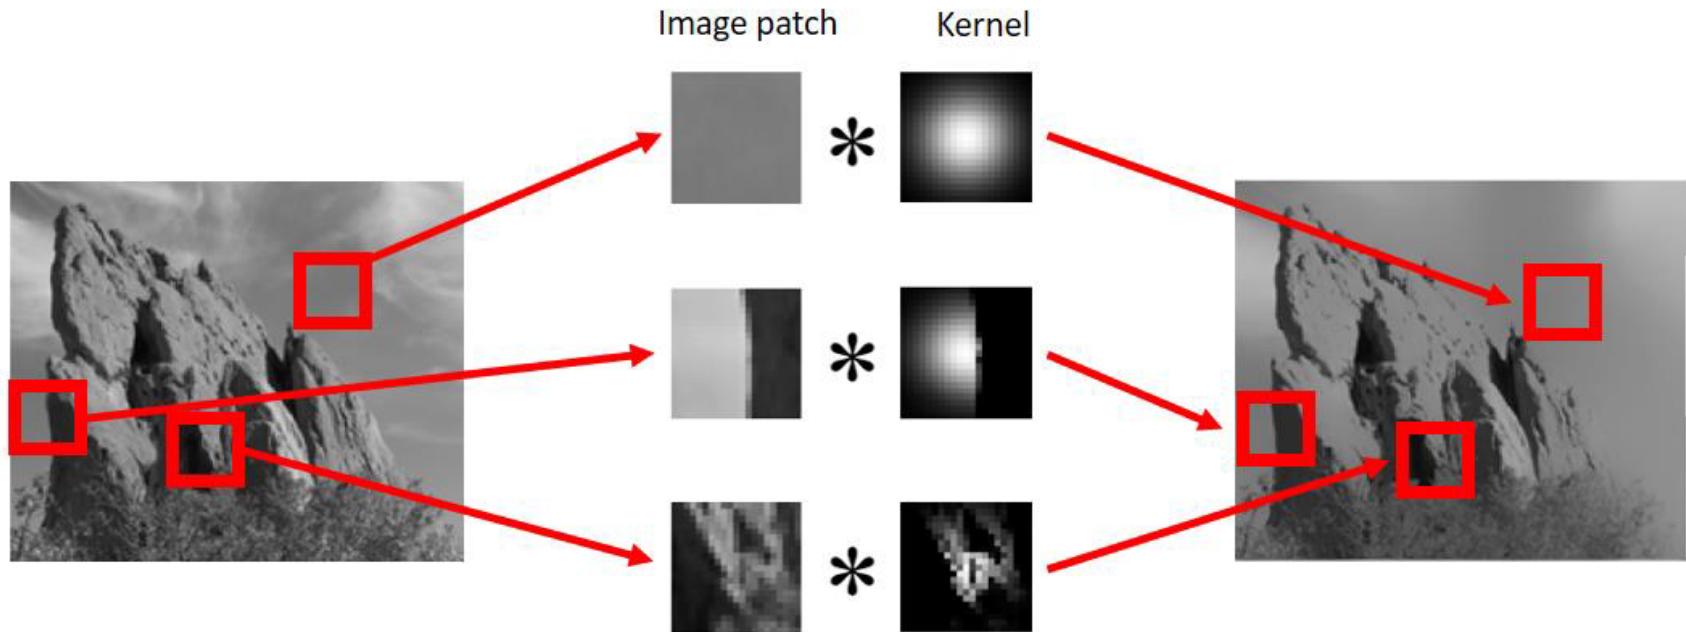
Bilateral filter

- **Gaussian filters do not preserve strong edges (discontinuities).** This is because they apply the same kernel everywhere.
- **Median filters** do preserve strong edges but remove small (weak) edges.



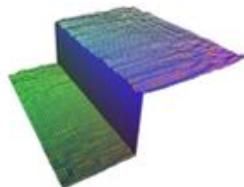
Bilateral filter

- **Gaussian filters do not preserve strong edges (discontinuities).** This is because they apply the same kernel everywhere.
- **Median filters** do preserve strong edges but remove small (weak) edges.
- **Bilateral filters** solve this by adapting the kernel locally to the intensity profile, so they are patch-content dependent: they only average pixels with similar brightness. Basically, the weighted average discards the influence of pixels with different brightness (across the discontinuity), while all the pixels that are on the same size of the discontinuity are smoothed.



Bilateral filter

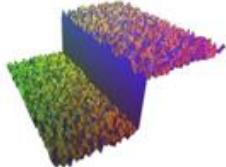
$$I'[x, y] = \frac{1}{W_p[x, y]} \sum_{u=-k}^k \sum_{v=-k}^k I[x - u, y - v] G_{\sigma_r}(I[x - u, y - v] - I[x, y]) G_{\sigma_s}[u, v]$$



smoothed image

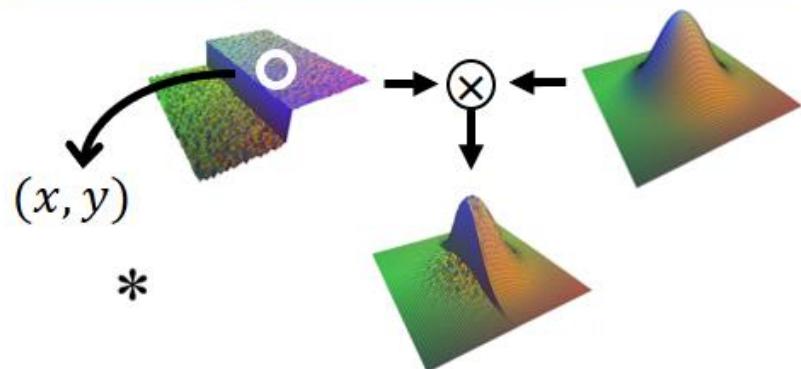
$$= \frac{1}{W_p[x, y]} \cdot$$

normalization
factor



input image

1-D Gaussian
range component spatial
component



$$W_p[x, y] = \sum_{u=-k}^k \sum_{v=-k}^k G_{\sigma_r}(I[x - u, y - v] - I[x, y]) G_{\sigma_s}[u, v]$$

Normalization factor
(so that the filter values sum up to 1)

Bilateral filter



input

larger neighborhoods
are smoothed



$\sigma_s = 18$



$\sigma_s = 6$



$\sigma_s = 2$

$\sigma_r = 0.1$



Stronger edges are
smoothed

$\sigma_r = 0.25$



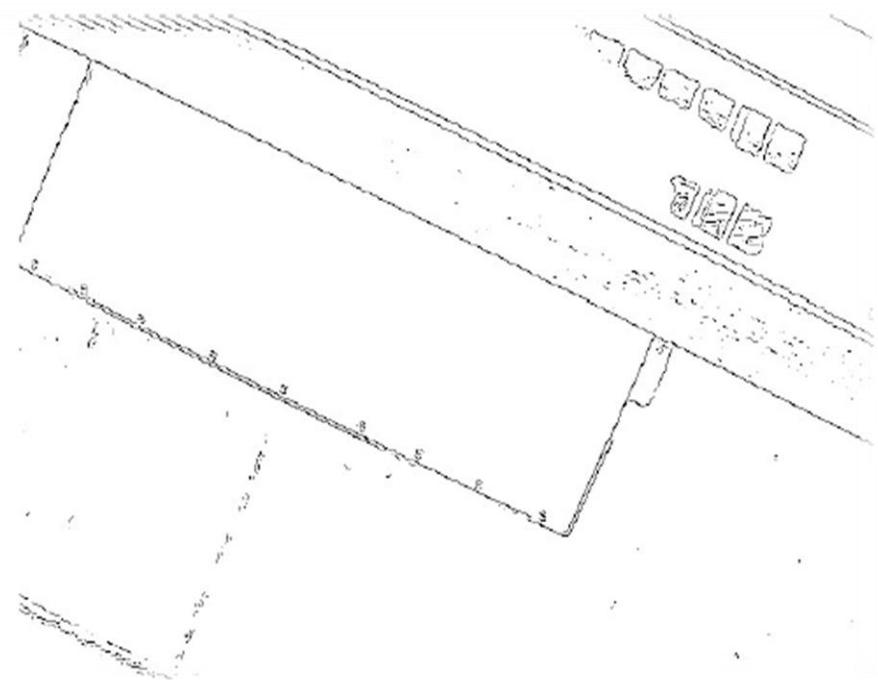
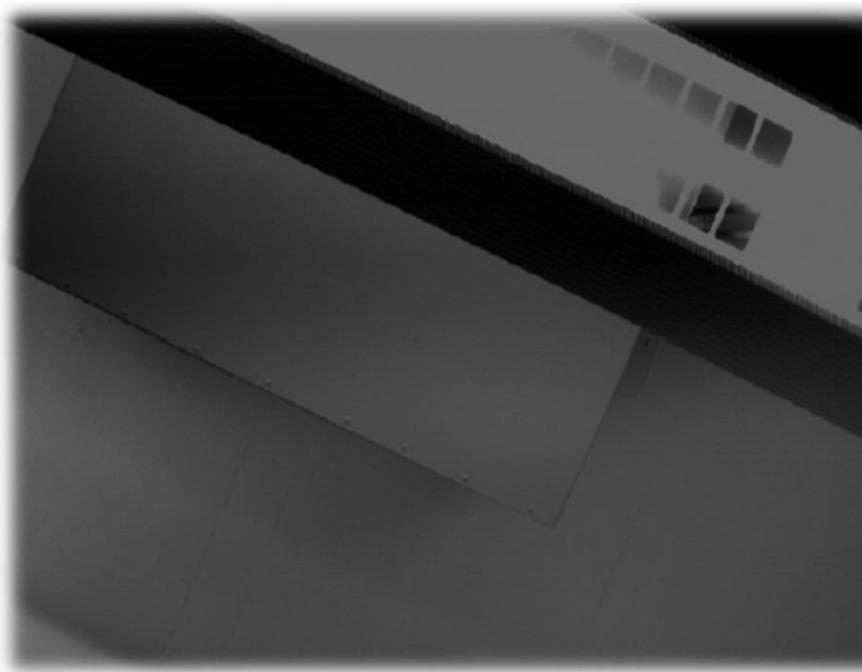
$\sigma_r = \infty$
(Gaussian blur)



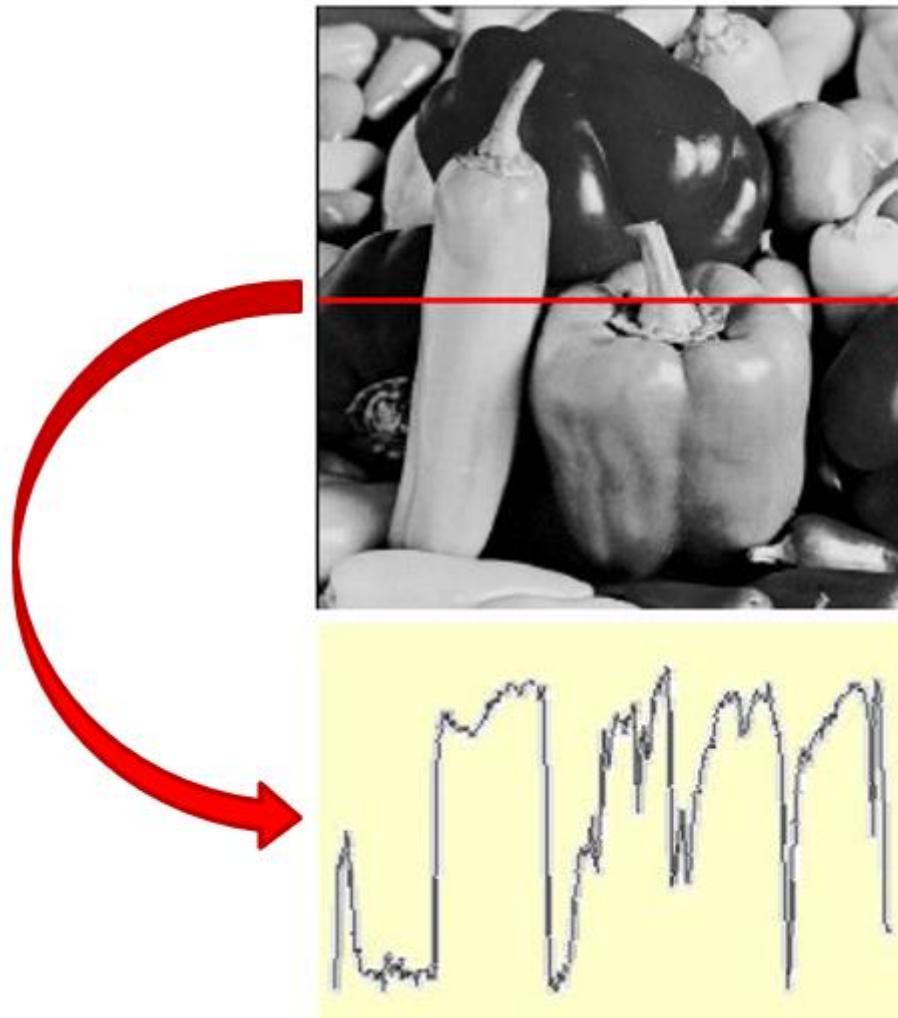
High-pass filtering (edge detection)

Edge detection

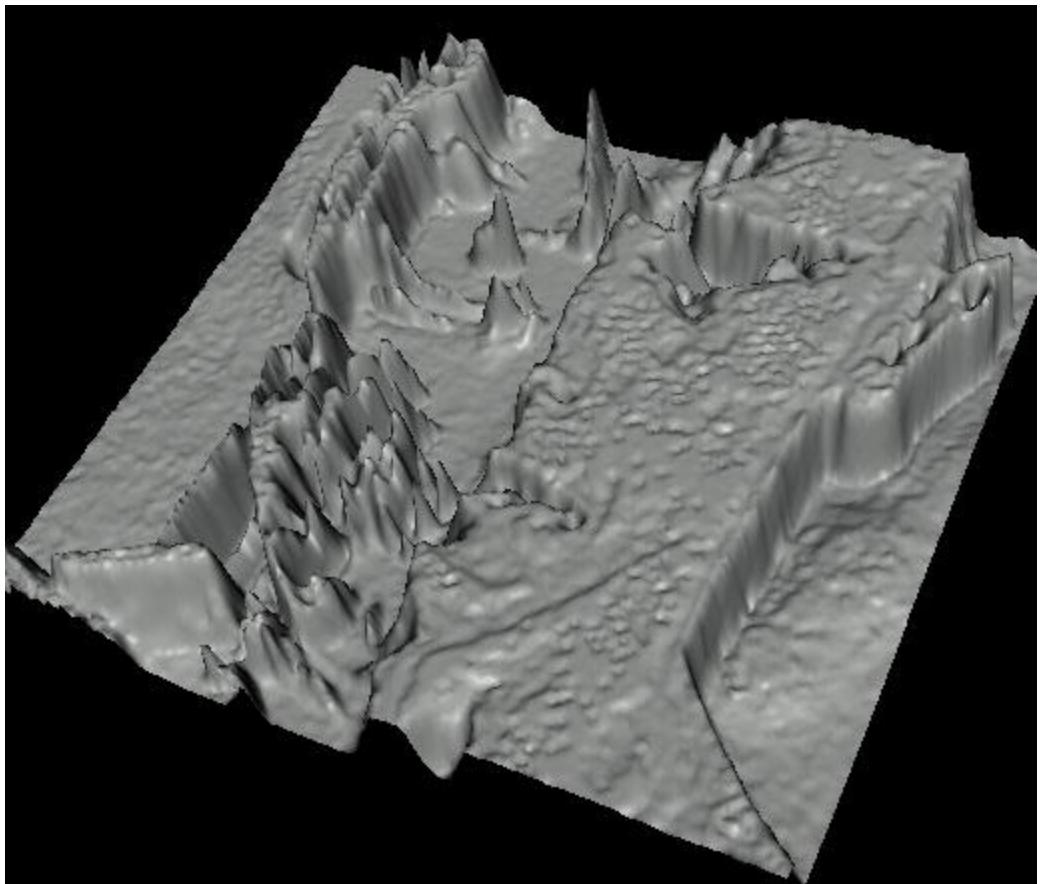
- Ultimate goal of edge detection: an idealized line drawing.
- Edge contours in the image correspond to important scene contours.



Edges are sharp intensity changes



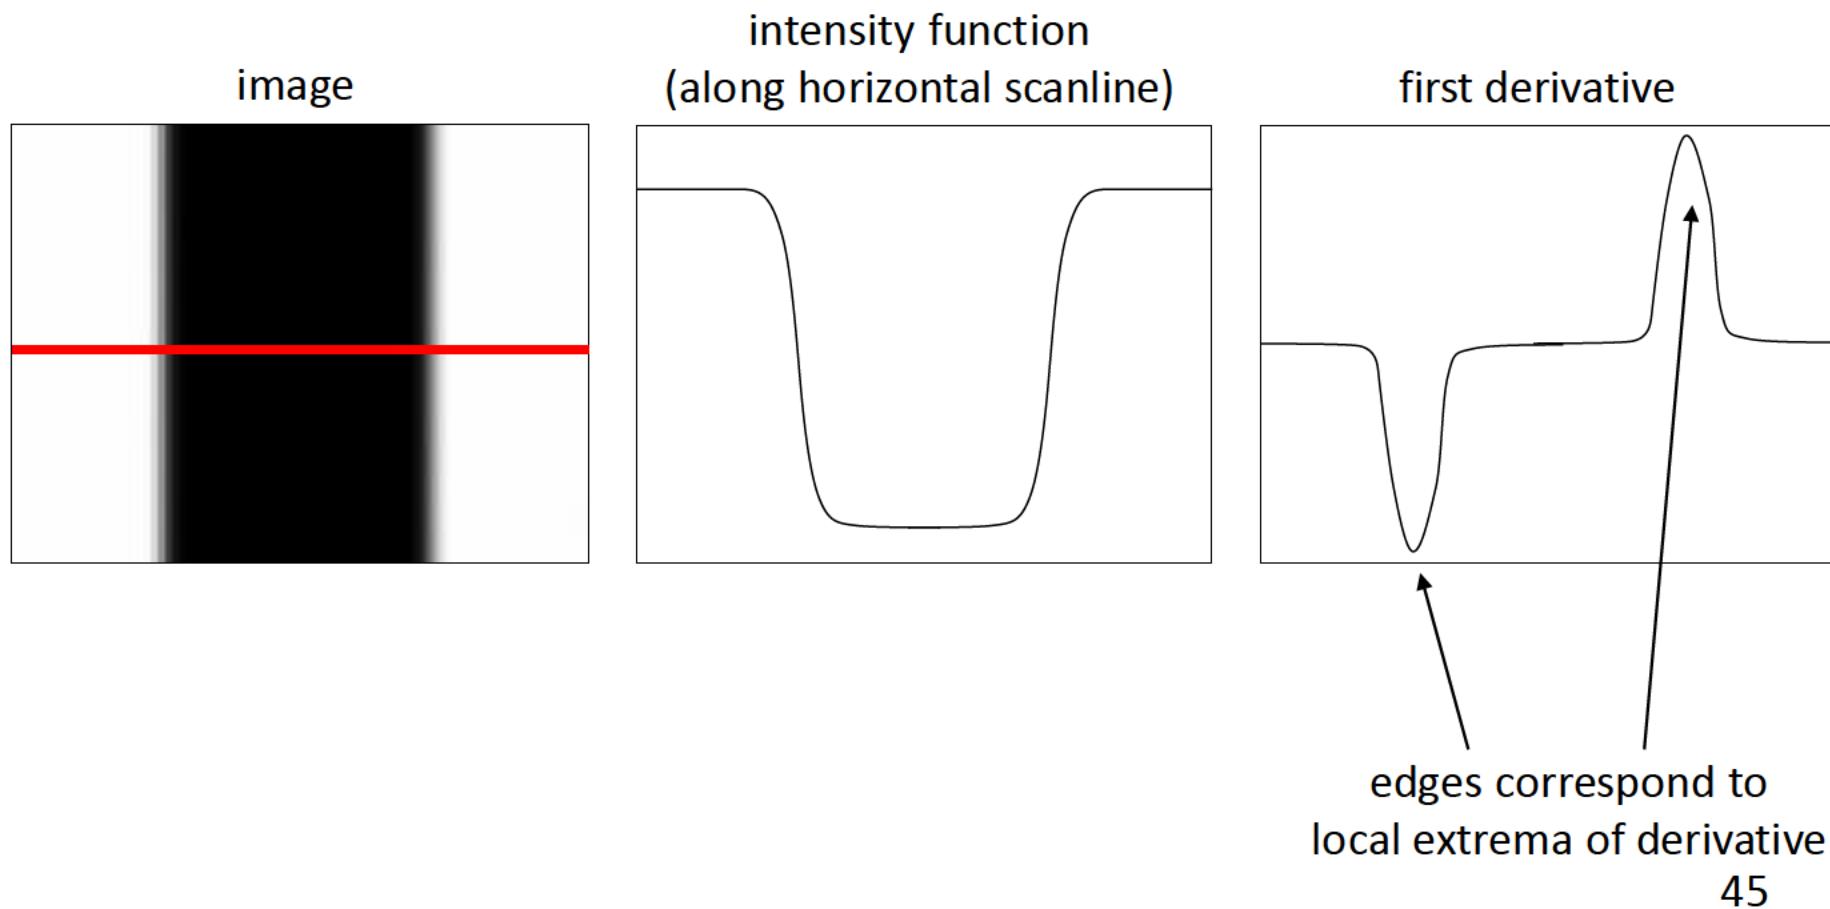
Images as functions $F(x, y)$



Edges look like steep cliffs

Derivatives and edges

An edge is a place of rapid change in the image intensity function.



Differentiation and convolution

For a 2D function $I(x, y)$ the partial derivative along x is:

$$\frac{\partial I(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{I(x + \varepsilon, y) - I(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

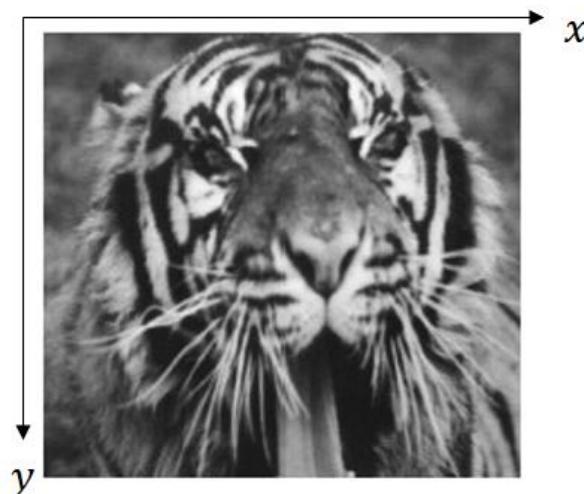
$$\frac{\partial I(x, y)}{\partial x} \approx \frac{I(x + 1, y) - I(x, y)}{1}$$

What would be the respective filters along x and y to implement the partial derivatives as a convolution?

Partial derivatives of an image

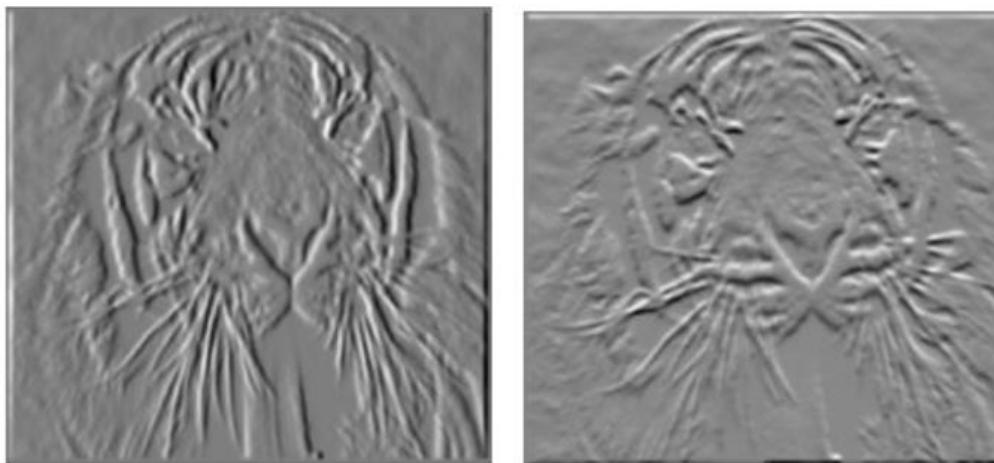
$$\frac{\partial I(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial I(x, y)}{\partial y}$$

-1
1



Alternative Finite-difference filters

Prewitt filter $G_x = \begin{bmatrix} -1 & 0 & +1 \\ -1 & 0 & +1 \\ -1 & 0 & +1 \end{bmatrix}$ and $G_y = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ +1 & +1 & +1 \end{bmatrix}$

Sobel filter $G_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix}$ and $G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix}$

Sample Matlab code

```
>> im = imread('lion.jpg');  
>> h = fspecial('sobel');  
>> outim = imfilter(double(im), h);  
>> imagesc(outim);  
>> colormap gray;
```

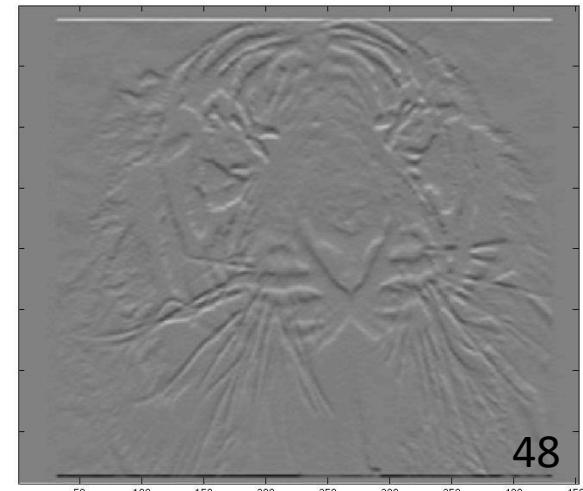
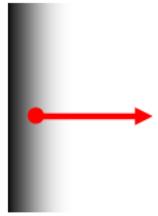


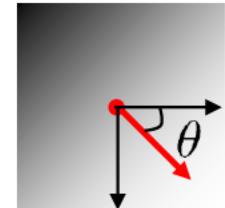
Image gradient

The gradient of an image: $\nabla I = \left[\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right]$

The gradient points in the direction of fastest intensity change


$$\nabla I = \left[\frac{\partial I}{\partial x}, 0 \right]$$


$$\nabla I = \left[0, \frac{\partial I}{\partial y} \right]$$


$$\nabla I = \left[\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right]$$

The gradient direction (orientation of edge normal) is given by:

$$\theta = \text{atan2} \left(\frac{\partial I}{\partial y}, \frac{\partial I}{\partial x} \right)$$

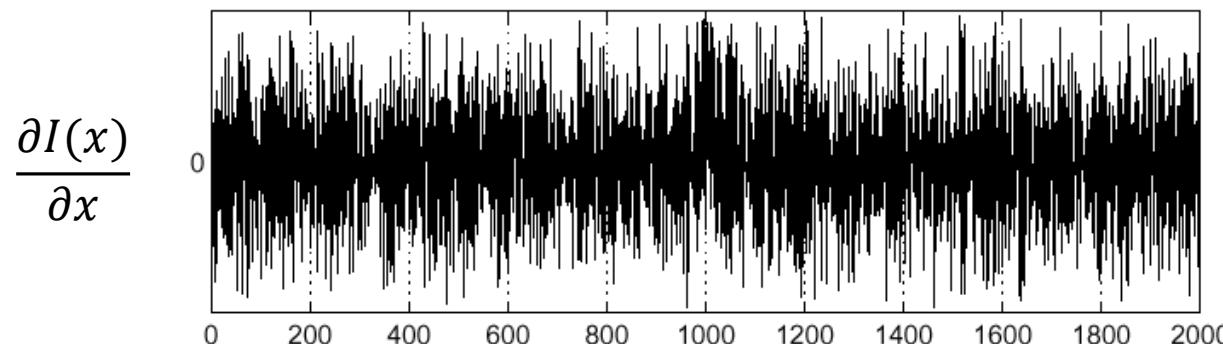
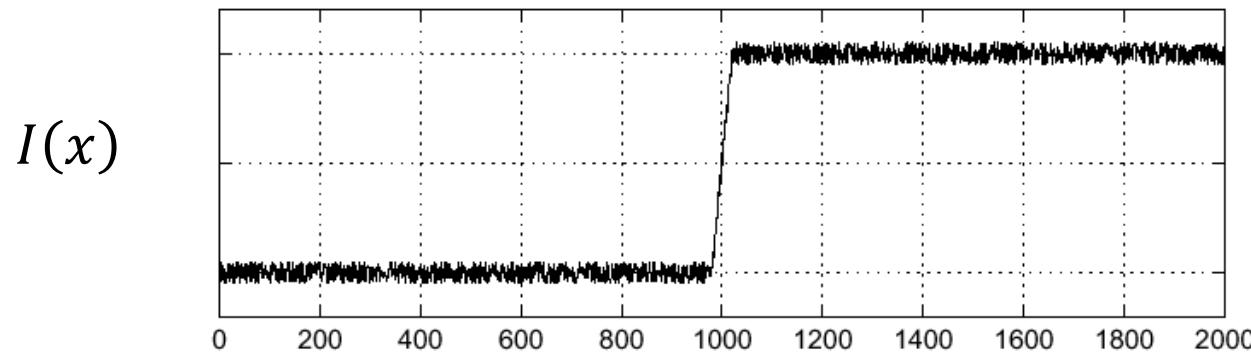
The *edge strength* is given by the gradient magnitude

$$\|\nabla I\| = \sqrt{\left(\frac{\partial I}{\partial x} \right)^2 + \left(\frac{\partial I}{\partial y} \right)^2}$$

Effects of noise

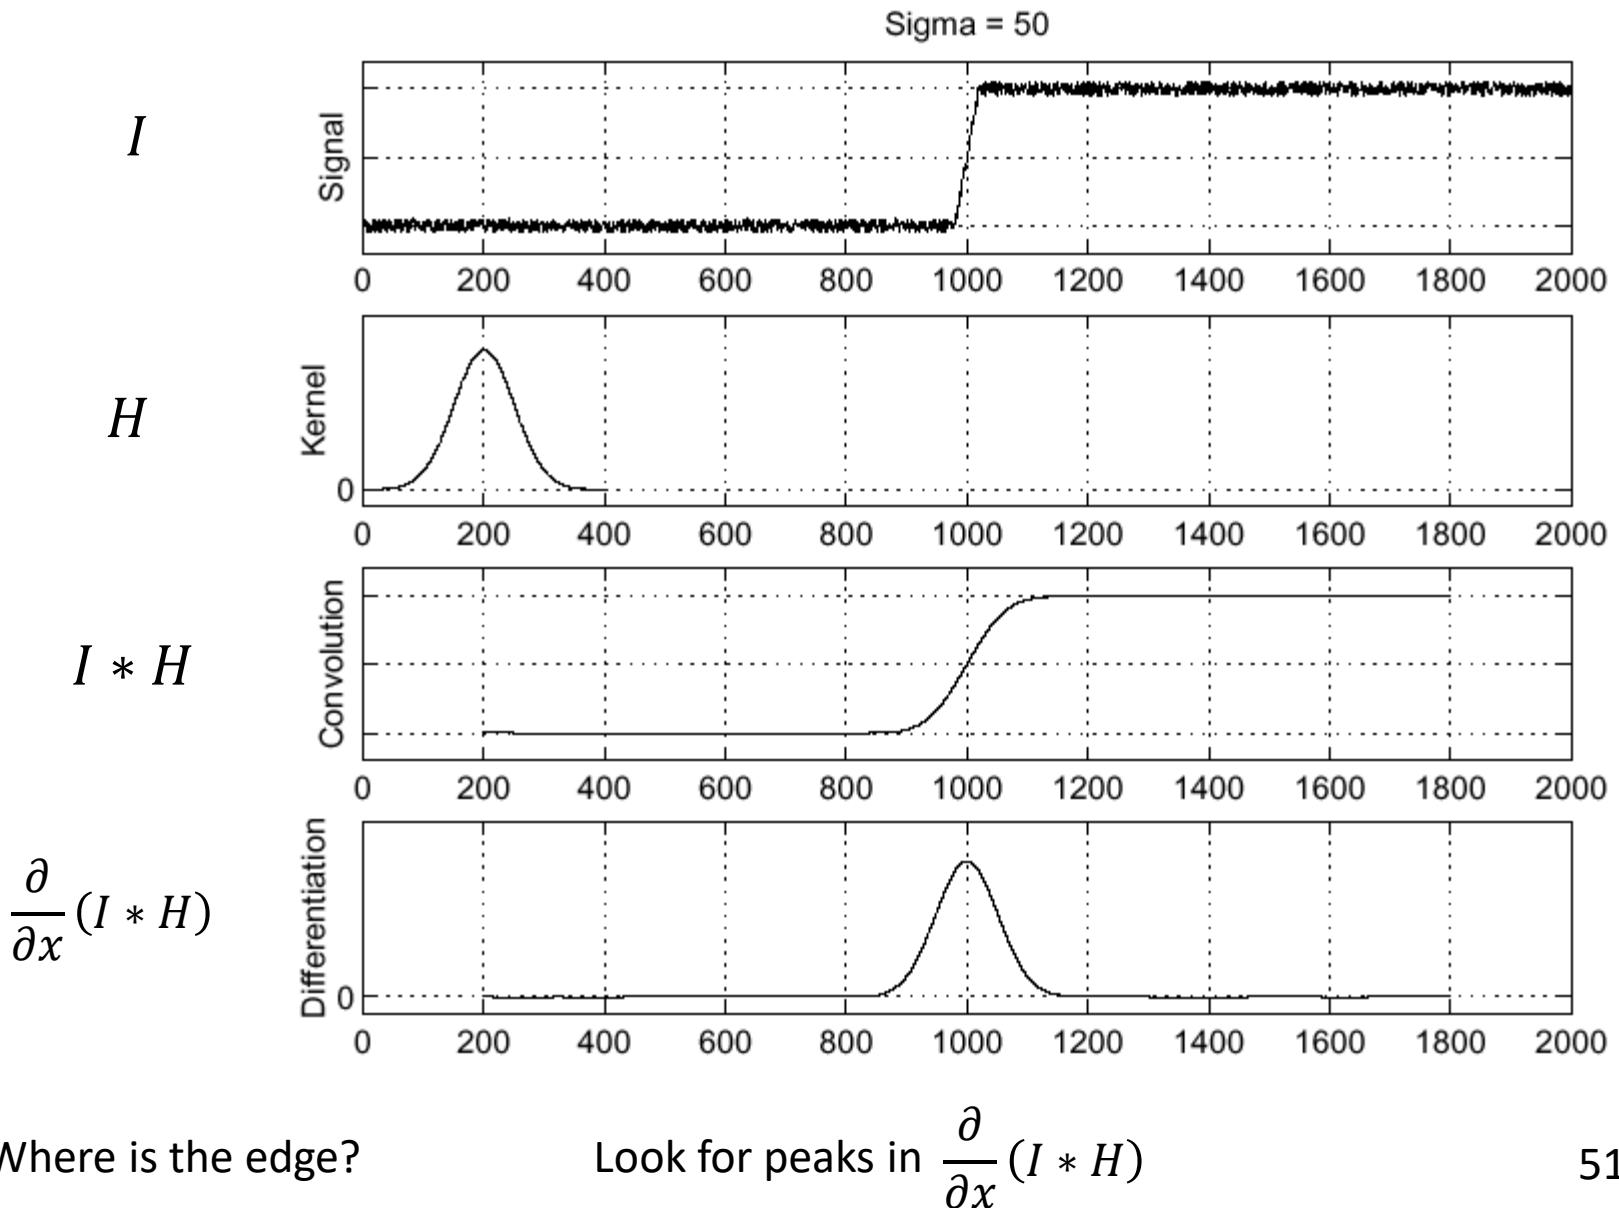
Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal



Where is the edge?

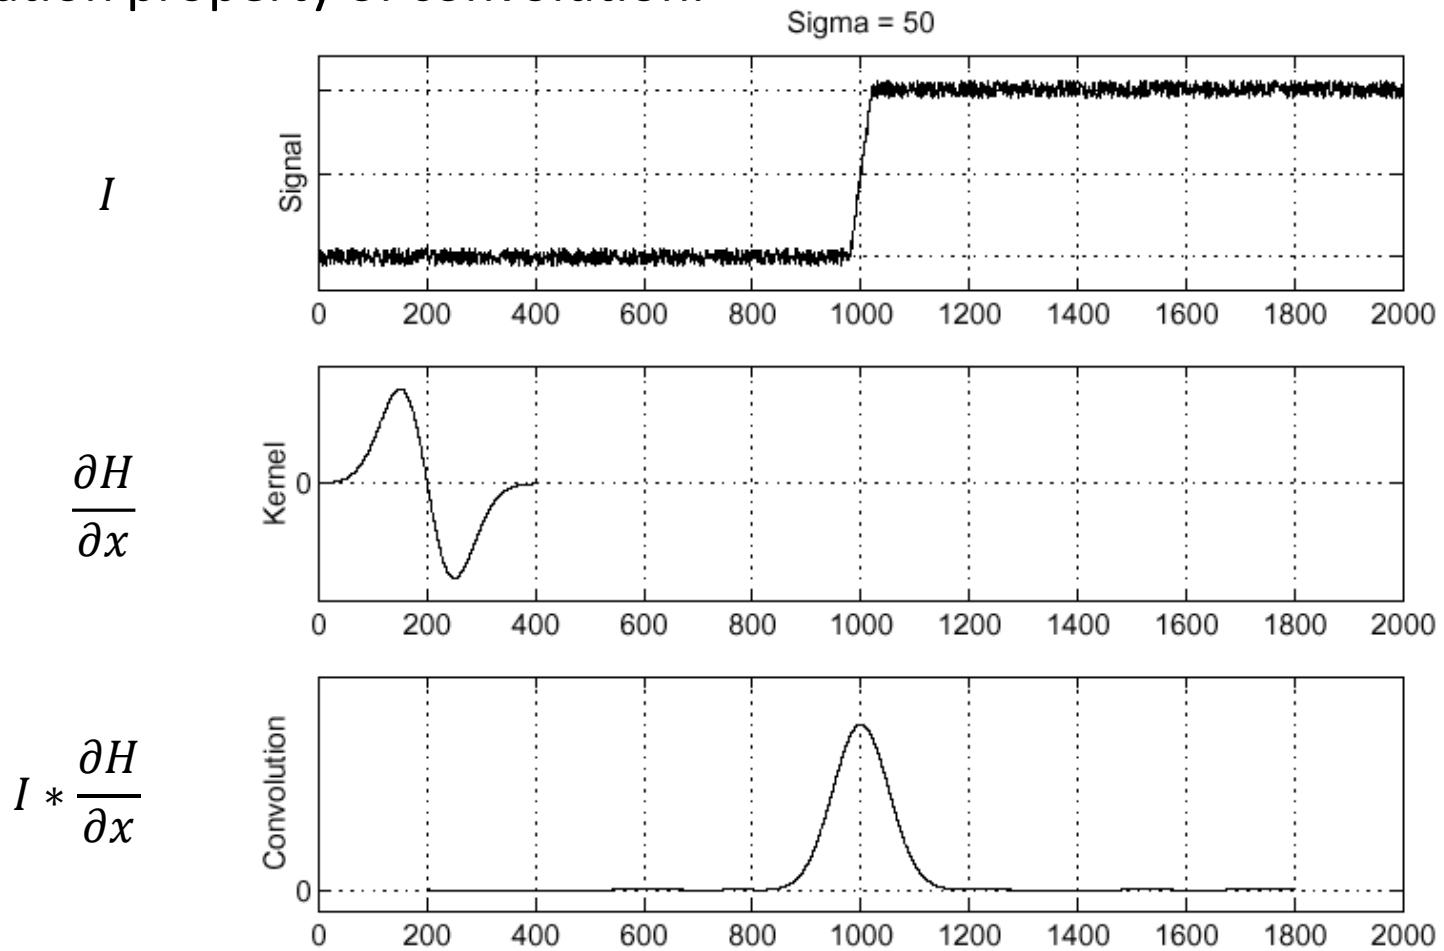
Solution: smooth first



Alternative: combine derivative and smoothing filter

$$\frac{\partial}{\partial x} (I * H) = I * \frac{\partial H}{\partial x}$$

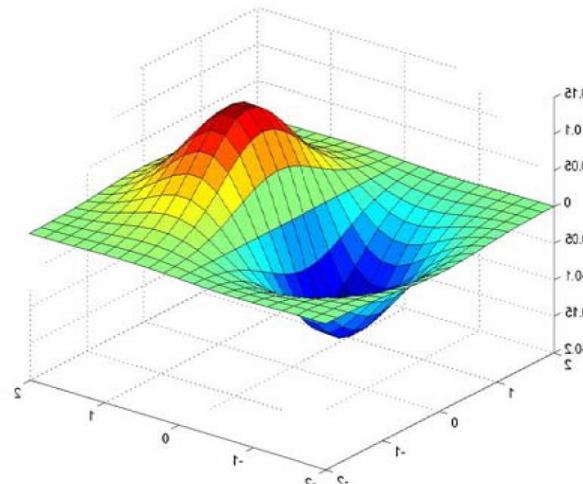
Differentiation property of convolution.



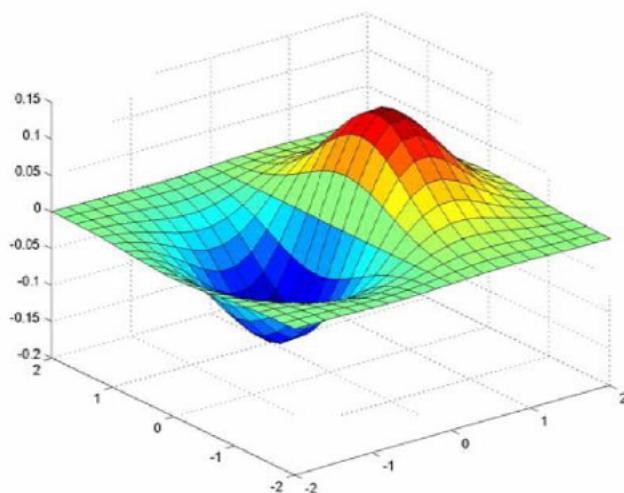
Derivative of Gaussian filter (along x)

$$(I * G) * H = I * (G * H)$$

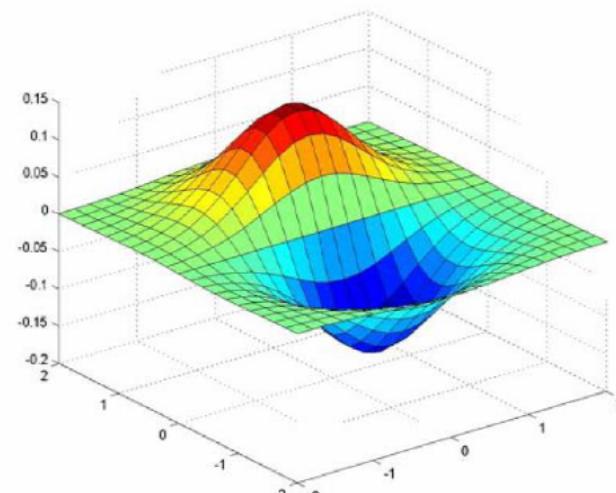
$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix} * \begin{bmatrix} 1 & -1 \end{bmatrix}$$



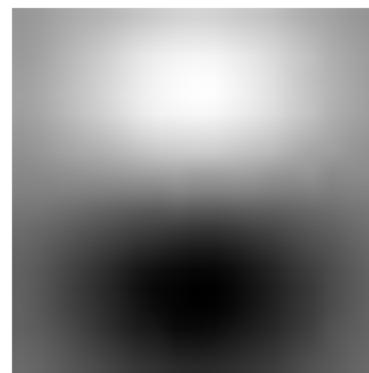
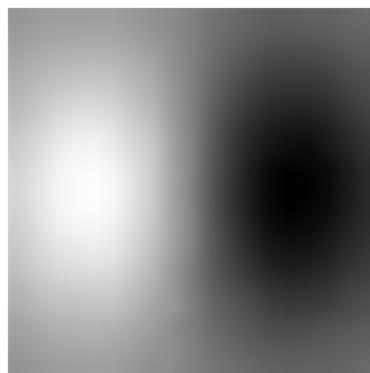
Derivative of Gaussian filters



x -direction

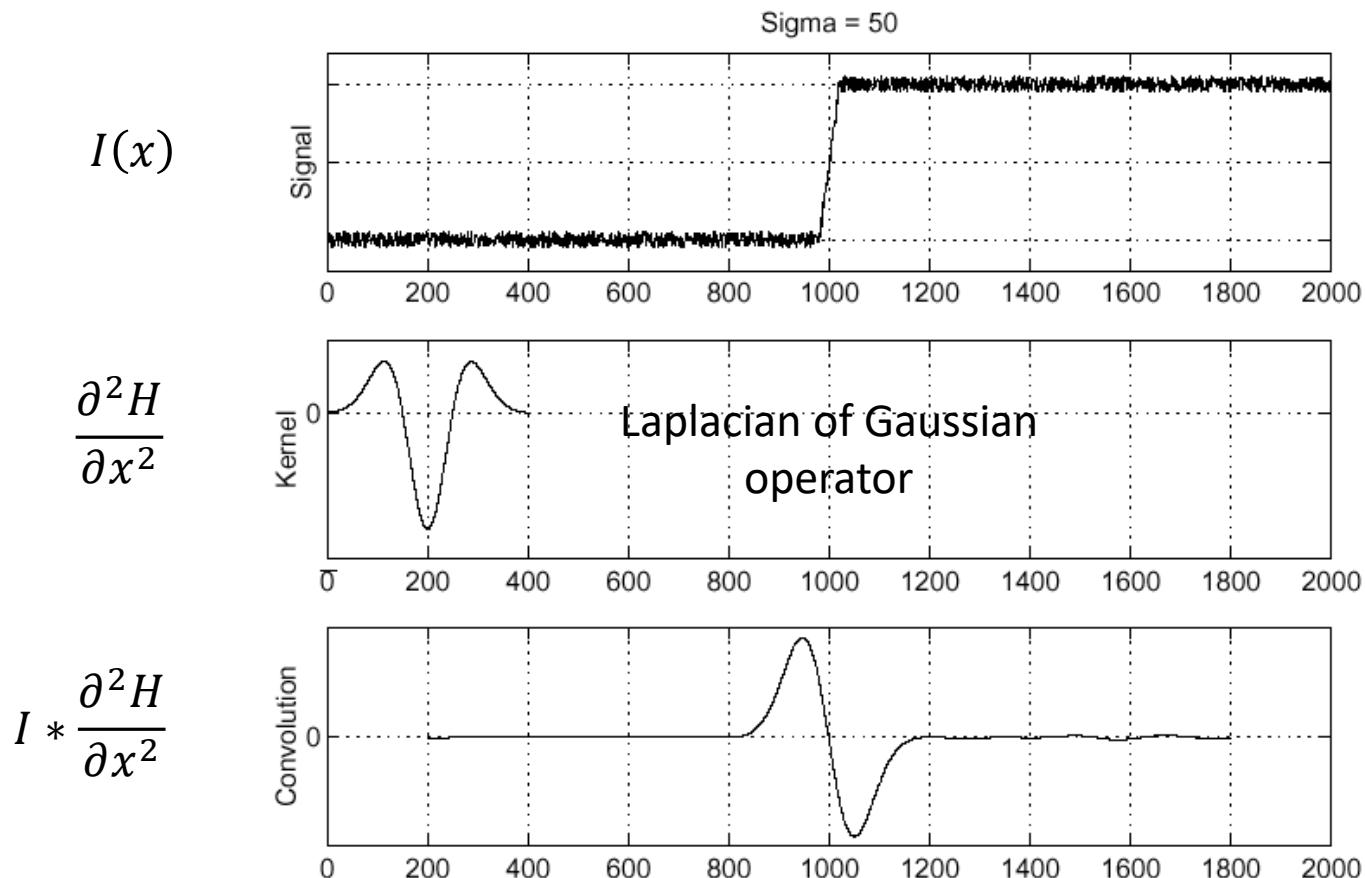


y -direction



Laplacian of Gaussian

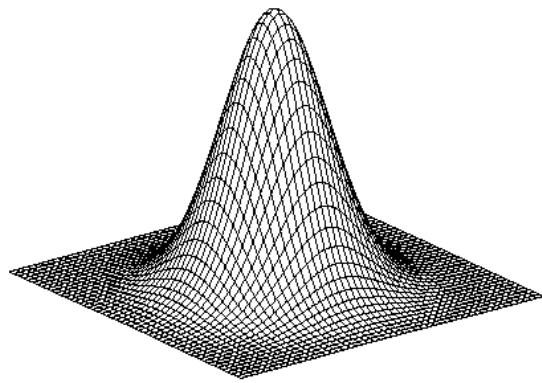
Consider $\frac{\partial^2}{\partial x^2}(I * H) = I * \frac{\partial^2 H}{\partial x^2}$



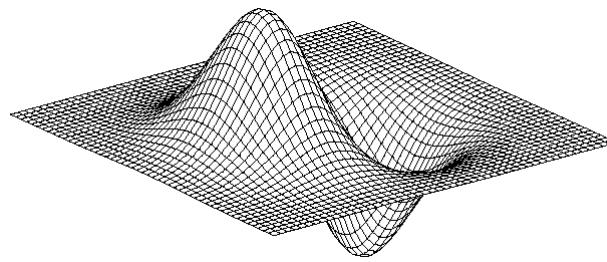
Where is the edge?

Zero-crossings of bottom graph

2D edge detection filters



Gaussian

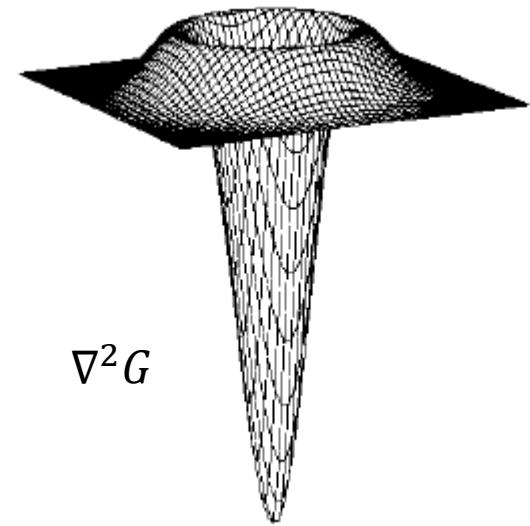


derivative of Gaussian

$$G = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$\frac{\partial G}{\partial x}$$

Laplacian of Gaussian



$$\nabla^2 G$$

- ∇^2 is the Laplacian operator: $\nabla^2 = \frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2}$

Summary on (linear) filters

- Smoothing filter:
 - has positive values
 - sums to 1 → preserve brightness of constant regions
 - removes “high-frequency” components: “low-pass” filter
- Derivative filter:
 - has opposite signs used to get high response in regions of high contrast
 - sums to 0 → no response in constant regions
 - highlights “high-frequency” components: “high-pass” filter

The Canny edge-detection algorithm (1986)



Original image (Lenna image: <https://en.wikipedia.org/wiki/Lenna>)

Take a grayscale image.
If not grayscale (i.g.,
RGB), convert it into a
grayscale by replacing
each pixel by the mean
value of its R, G, B
components.

The Canny edge-detection algorithm (1986)



Take a grayscale image.
If not grayscale (i.g.,
RGB), convert it into a
grayscale by replacing
each pixel by the mean
value of its R, G, B
components.

Original image (Lenna image: <https://en.wikipedia.org/wiki/Lenna>)

The Canny edge-detection algorithm (1986)

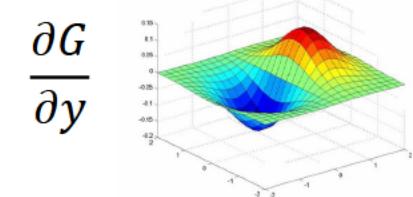
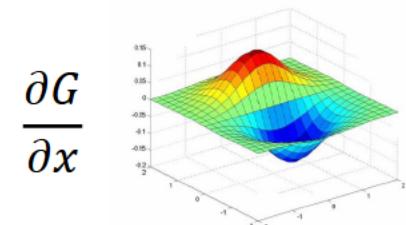


$$\|\nabla I\| = \sqrt{\left(\frac{\partial I}{\partial x}\right)^2 + \left(\frac{\partial I}{\partial y}\right)^2} : \text{Edge strength}$$

Convolve the image I with x and y derivatives of Gaussian filter

$$\frac{\partial I}{\partial x} = I * \frac{\partial G}{\partial x}$$

$$\frac{\partial I}{\partial y} = I * \frac{\partial G}{\partial y}$$



The Canny edge-detection algorithm (1986)



Thresholded $\|\nabla I\|$

Threshold it (i.e., set to 0 all pixels whose value is below a given threshold)

The Canny edge-detection algorithm (1986)



Take local maximum
along gradient direction

Thinning: non-maxima suppression (local-maxima detection)
along edge direction

Summary

- Image filtering (definition, motivation, applications)
- Moving average
- Linear filters and formulation: box filter, Gaussian filter
- Boundary issues
- Non-linear filters
 - Median & bilateral filters
- Edge detection
 - Derivating filters (Prewitt, Sobel)
 - Combined derivative and smoothing filters (deriv. of Gaussian)
 - Laplacian of Gaussian
 - Canny edge detector
- Readings: Ch. 3.2, 4.2.1 of Szeliski book